

Machine Learning



Machine Learning for Bitcoin price prediction

Bitcoin price prediction

A comprehensive analysis using Linear Regression, Random Forest, and Gradient Boosting models

Natasha Lertsansiri 1650900234

Thanathon Satthayaphan 1650900234

Bitcoin price prediction



Overview :

- Introduction
- Data Overview
- Data Preprocessing
- Model Selection
- Model Training
- Testing and Evaluation
- Results
- Conclusion

Introduction

- We delve into the exciting realm of Bitcoin prediction through machine learning. Bitcoin, the pioneering cryptocurrency, has captivated the world with its meteoric rise and volatility. Understanding its price dynamics is not only a matter of financial interest but also a significant challenge due to its complex and often unpredictable nature.
- In this project, we harness the power of machine learning algorithms to forecast Bitcoin prices. By analyzing historical data, identifying patterns, and leveraging advanced predictive models, we aim to provide valuable insights into the future movements of this digital asset.

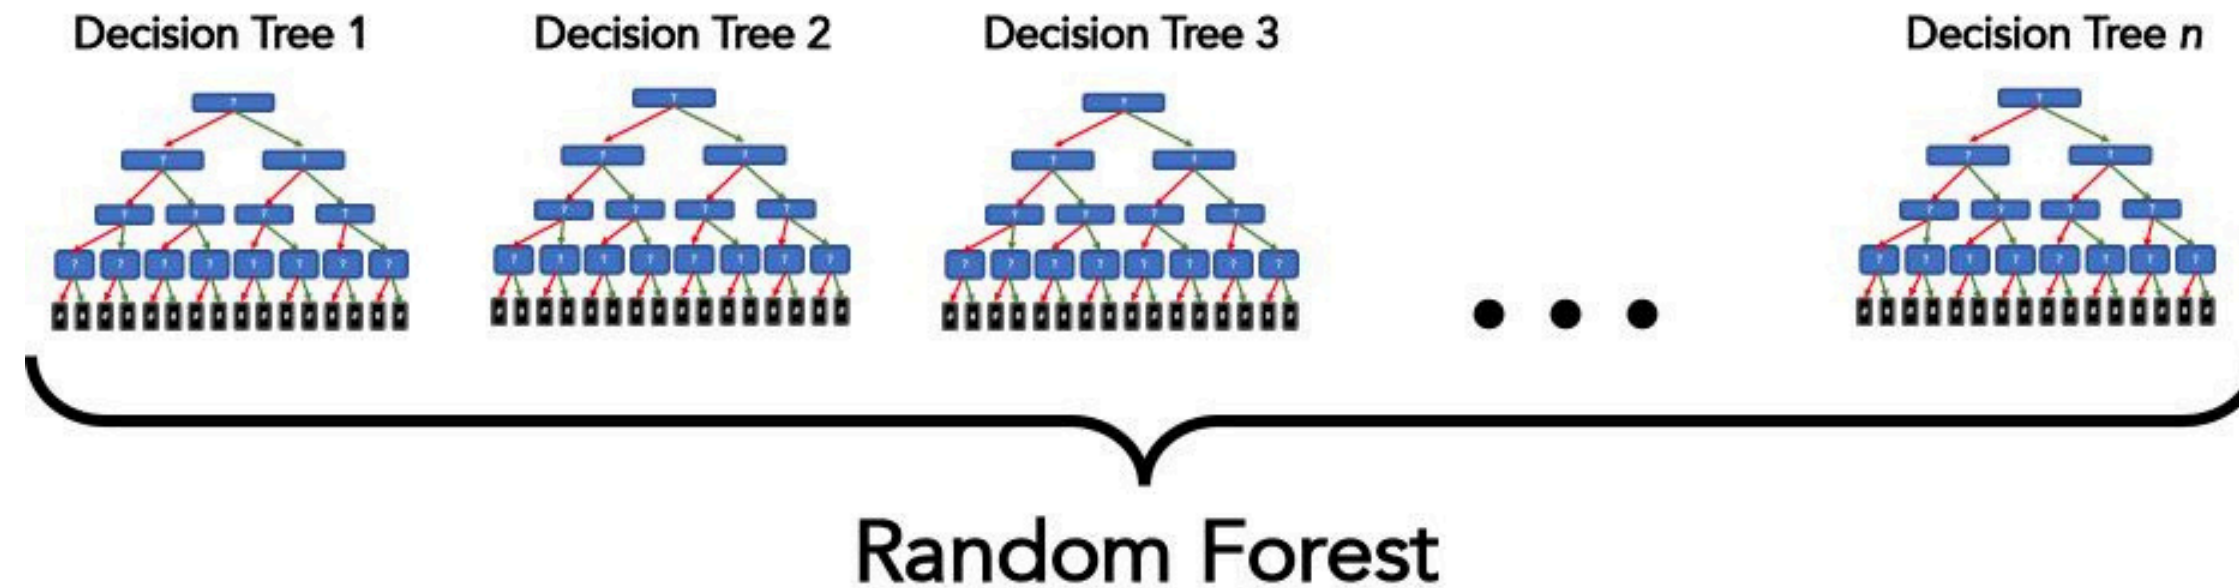
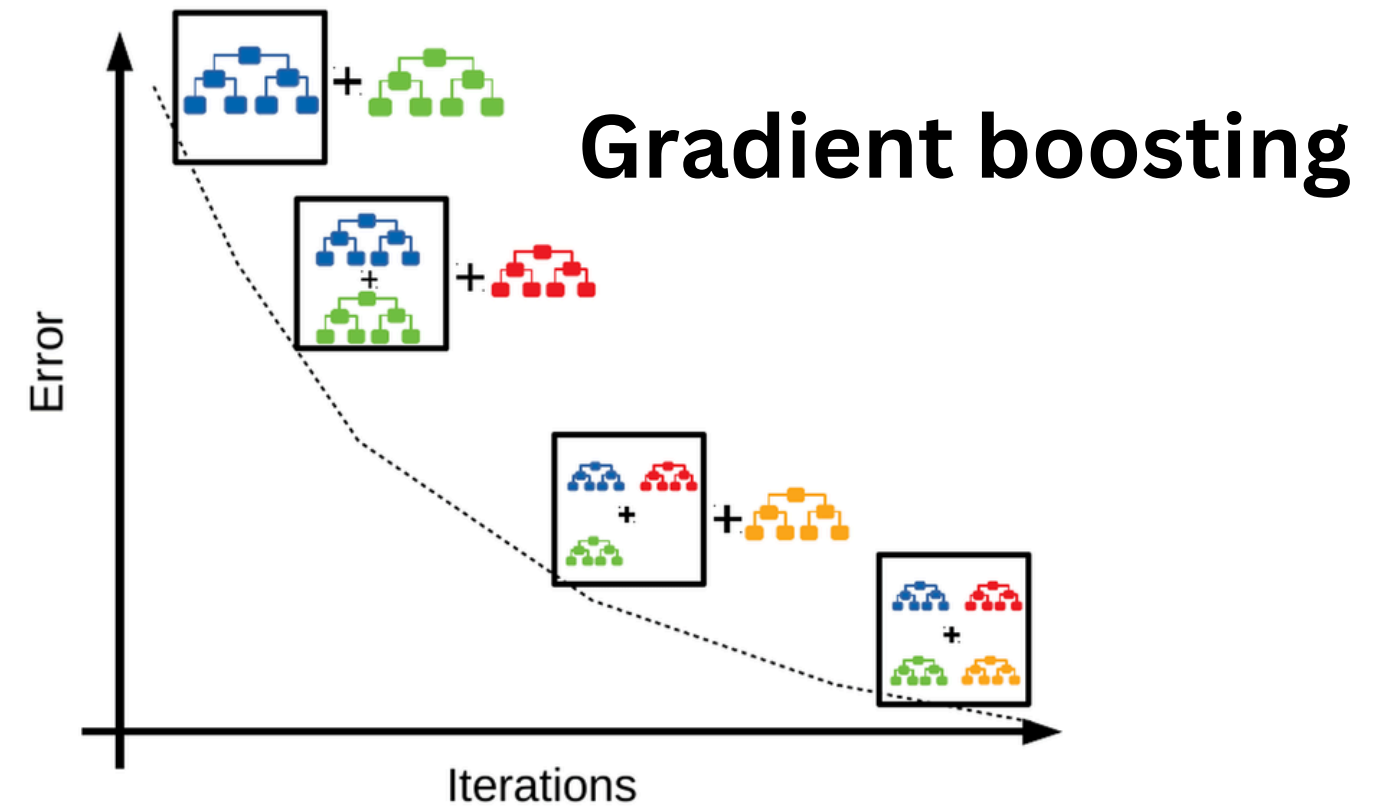

Data Overview

- **Data Source:**
<https://coincodex.com/crypto/bitcoin/historical-data/>
- **Key Features:** List the main features used (Open, High, Low, Volume, Market Cap).
- **Date Range:** Mention the time period the data covers.

Start	End	Open	High	Low	Close	Volume	Market Cap
1/1/2011	1/2/2011	0.3	0.3	0.3	0.3	0	0
1/2/2011	1/3/2011	0.3	0.3	0.3	0.3	0	0
1/3/2011	1/4/2011	0.295	0.295	0.295	0.295	0	0
1/4/2011	1/5/2011	0.299	0.299	0.299	0.299	0	0
1/5/2011	1/6/2011	0.299	0.299	0.299	0.299	0	0
1/6/2011	1/7/2011	0.298	0.298	0.298	0.298	0	0
1/7/2011	1/8/2011	0.32	0.32	0.32	0.32	0	0
1/8/2011	1/9/2011	0.3229	0.3229	0.3229	0.3229	0	0
1/9/2011	1/10/2011	0.323	0.323	0.323	0.323	0	0
1/10/2011	1/11/2011	0.3266	0.3266	0.3266	0.3266	0	0
1/11/2011	1/12/2011	0.3266	0.3266	0.3266	0.3266	0	0
1/12/2011	1/13/2011	0.3188	0.3188	0.3188	0.3188	0	0
1/13/2011	1/14/2011	0.3176	0.3176	0.3176	0.3176	0	0
1/14/2011	1/15/2011	0.4	0.4	0.4	0.4	0	0
1/15/2011	1/16/2011	0.386	0.386	0.386	0.386	0	0
1/16/2011	1/17/2011	0.3868	0.3868	0.3868	0.3868	0	0
1/17/2011	1/18/2011	0.3495	0.3495	0.3495	0.3495	0	0
1/18/2011	1/19/2011	0.313	0.313	0.313	0.313	0	0
1/19/2011	1/20/2011	0.313	0.313	0.313	0.313	0	0
1/20/2011	1/21/2011	0.39	0.39	0.39	0.39	0	0
1/21/2011	1/22/2011	0.4199	0.4199	0.4199	0.4199	0	0
1/22/2011	1/23/2011	0.4443	0.4443	0.4443	0.4443	0	0
1/23/2011	1/24/2011	0.4424	0.4424	0.4424	0.4424	0	0
1/24/2011	1/25/2011	0.4199	0.4199	0.4199	0.4199	0	0
1/25/2011	1/26/2011	0.41	0.41	0.41	0.41	0	0
1/26/2011	1/27/2011	0.417	0.417	0.417	0.417	0	0
1/27/2011	1/28/2011	0.4212	0.4212	0.4212	0.4212	0	0
1/28/2011	1/29/2011	0.446	0.446	0.446	0.446	0	0
1/29/2011	1/30/2011	0.439	0.439	0.439	0.439	0	0

Model Selection

Linear Regression



Libraries and Tools used

- Pandas
- sklearn
- numpy
- matplotlib
- Visual Studio Code - IDE

Data Preprocessing

Missing data and Convert to datetime

```
1 import pandas as pd
2
3 # โหลดข้อมูล
4 data = pd.read_csv('bitcoin_dialy.csv')
5 print(data.head())
6
7 # ตรวจสอบmissing data
8 missing_data = data.isnull().sum()
9
10 # แปลงคอลัมน์ 'Start' และ 'End' ให้เป็น datetime
11 data['Start'] = pd.to_datetime(data['Start'])
12 data['End'] = pd.to_datetime(data['End'])
13
14 print("Missing data:\n", missing_data)
15 print("\nData types after conversion:\n", data.dtypes)
16
17 from sklearn.model_selection import train_test_split
18 # ตัวแปรอิสระ (features)
19 X = data[['Open', 'High', 'Low', 'Volume', 'Market Cap']]
20 # ตัวแปรตาม (target)
21 y = data['Close']
22 # แบ่งข้อมูลtrain/test 80/20
23 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
24 # ขนาดของชุดข้อมูลที่แบ่งได้
25 print("Training set:", X_train.shape, y_train.shape)
26 print("Test set:", X_test.shape, y_test.shape)
```

	Start	End	Open	High	Low	Close	Volume	Market Cap
0	2011-01-01	2011-01-02	0.300	0.300	0.300	0.300	0.0	0.0
1	2011-01-02	2011-01-03	0.300	0.300	0.300	0.300	0.0	0.0
2	2011-01-03	2011-01-04	0.295	0.295	0.295	0.295	0.0	0.0
3	2011-01-04	2011-01-05	0.299	0.299	0.299	0.299	0.0	0.0
4	2011-01-05	2011-01-06	0.299	0.299	0.299	0.299	0.0	0.0

Missing data:

Start	0
End	0
Open	0
High	0
Low	0
Close	0
Volume	0
Market Cap	0

dtype: int64

Data types after conversion:

Start	datetime64[ns]
End	datetime64[ns]
Open	float64
High	float64
Low	float64
Close	float64
Volume	float64
Market Cap	float64

dtype: object

Training set: (3799, 5) (3799,)
Test set: (950, 5) (950,)

Model Training and Evaluation

```
28 from sklearn.linear_model import LinearRegression
29 from sklearn.metrics import mean_squared_error, r2_score
30 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
31
32 # สร้างและ train Linear Regression model
33 linear_model = LinearRegression()
34 linear_model.fit(X_train, y_train)
35
36 # Predict test data
37 y_pred_linear = linear_model.predict(X_test)
38
39 # ประเมิน Linear Regression model
40 mse_linear = mean_squared_error(y_test, y_pred_linear)
41 r2_linear = r2_score(y_test, y_pred_linear)
42
43 # สร้างและ train Random Forest Regression model
44 random_forest_model = RandomForestRegressor(random_state=42)
45 random_forest_model.fit(X_train, y_train)
46 y_pred_rf = random_forest_model.predict(X_test)
47
48 # ประเมิน Random Forest Regression model
49 mse_rf = mean_squared_error(y_test, y_pred_rf)
50 r2_rf = r2_score(y_test, y_pred_rf)
51
52 # สร้างและ train Gradient Boosting Regression model
53 gb_model = GradientBoostingRegressor(random_state=42)
54 gb_model.fit(X_train, y_train)
55 y_pred_gb = gb_model.predict(X_test)
56
57 # ประเมิน
58 mse_gb = mean_squared_error(y_test, y_pred_gb)
59 r2_gb = r2_score(y_test, y_pred_gb)
60
61 model_evaluation_results = {
62     "Linear Regression": {"MSE": mse_linear, "R2": r2_linear},
63     "Random Forest Regression": {"MSE": mse_rf, "R2": r2_rf},
64     "Gradient Boosting Regression": {"MSE": mse_gb, "R2": r2_gb}
65 }
66 model_evaluation_results
67
68 print("Linear Regression - MSE:", mse_linear, "\nR2 Score:", r2_linear)
69 print("Random Forest Regression - MSE:", mse_rf, "\nR2 Score:", r2_rf)
70 print("Gradient Boosting Regression - MSE:", mse_gb, "\nR2 Score:", r2_gb)
```

Linear Regression - MSE: 66069.4376948983 R2 Score: 0.9997035066528899

Random Forest Regression - MSE: 114626.86487298366 R2 Score: 0.999485600240888

Gradient Boosting Regression - MSE: 159301.79658384464 R2 Score: 0.999285116923684

**Linear Regression - MSE: 66069.4376948983
R2 Score: 0.9997035066528899**

**Random Forest Regression - MSE: 114626.86487298366
R2 Score: 0.999485600240888**


**Gradient Boosting Regression - MSE: 159301.79658384464
R2 Score: 0.999285116923684**

Model Prediction and Evaluation

```
75 data['Start'] = pd.to_datetime(data['Start'])
76 data.set_index('Start', inplace=True)
77
78 data['Time_Index'] = (data.index - data.index.min()) / pd.Timedelta(days=1)
79 X = data[['Time_Index']]
80 y = data['Close']
81
82 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
83
84 rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
85 rf_model.fit(X_train, y_train)
86
87 y_pred_rf = rf_model.predict(X_test)
88 print('\n\n', y_pred_rf)
89
90
91 plt.figure(figsize=(10, 5))
92 plt.plot(data.index, y, label='Actual Price', color='lime')
93 plt.scatter(X_test.index, y_pred_rf, color='deeppink', label='Predicted Price', s=10)
94 plt.title('Comparison of True Prices and Predicted Prices by Dates (Random Forest)')
95 plt.xlabel('Date')
96 plt.ylabel('Price')
97 plt.legend()
98 plt.grid(True)
99 plt.show()
```

```
4.30974910e+02 2.16539433e+04 1.65266350e+01 1.00392109e+04
6.64310867e+03 3.69073926e+04 9.02079700e+02 6.00140020e+02
2.30943483e+04 6.46728200e+02 2.41997206e+04 1.44235930e+03
1.10359310e+01 4.82246425e+03 2.97545019e+04 8.20734742e+03
3.88274606e+03 5.70189965e+04 2.01788699e+04 7.09975422e+03
7.44670000e-01 8.27310458e+03 4.24771442e+04 6.45015800e+00
2.69529310e+01 3.31814971e+04 1.05646519e+04 1.22475830e+03
1.23997400e+01 2.78136586e+04 4.85557295e+04 3.21745200e+00
6.26126400e+00 8.29160910e+03 4.56914016e+04 3.48964084e+04
8.20044171e+03 2.68071680e+02 7.49726220e+01 1.14610510e+02
1.70317986e+04 2.15056940e+01 1.13101840e+03 8.70363116e+03
4.02104180e+02 6.93310633e+03 1.24279770e+03 6.13981610e+02
3.87367700e+00 7.52004600e+00 6.41697000e+02 6.00925540e+02
4.57924000e+02 6.87017656e+03 9.19689838e+03 6.11632642e+03
3.91671800e+00 1.34815710e+03 7.39462592e+03 2.87601200e+00
5.03520891e+04 2.75781250e+03 7.28632908e+03 2.29509410e+02
8.55242600e+00 5.56261600e+00 3.86434680e+02 6.03853300e+00
1.36093880e+01 6.40164000e+00 6.42031714e+03 6.89498171e+03
2.08493362e+04 6.40493640e+02 1.67843015e+04 1.70176585e+04
3.54321247e+03 6.77225300e+03 1.35492335e+04 4.36578105e+04
3.91166250e+02 3.03332725e+04 8.22946020e+02 1.17507130e+01
9.28434482e+03 4.53737370e+01 2.35130410e+04 2.68139466e+04
4.30810000e-01 1.72164999e+04 4.23052058e+04 2.75633076e+04
4.14068320e+02 2.69019710e+04 9.73524785e+03 7.61485791e+03
2.30762150e+02 1.01531626e+04 4.23671930e+02 1.02235931e+04
9.40549866e+03 7.83147821e+03 8.76204741e+03 4.48864240e+03
4.45574260e+02 3.87758413e+04 6.30801080e+02 1.17130237e+04
1.19035550e+01 4.04687470e+02 7.32909137e+03 2.12647198e+04
3.71017770e+01 1.40900800e+02 6.63271500e+00 5.37035583e+03
8.88241157e+03 1.07406270e+03 3.68353636e+03 9.75675370e+01
```

Result




Model Prediction and Evaluation

```
101 start_date = '2018-01-01'
102 end_date = '2020-12-31'
103
104 data_filtered = data.loc[start_date:end_date]
105 X_filtered = X.loc[start_date:end_date]
106 y_filtered = y.loc[start_date:end_date]
107
108 y_pred_filtered = rf_model.predict(X_filtered)
109 print('\n\n',y_pred_filtered)
110
111 plt.figure(figsize=(10, 5))
112 plt.plot(data_filtered.index, y_filtered, label='Actual Price', color='lime')
113 plt.scatter(data_filtered.index, y_pred_filtered, color='deeppink', label='Predicted Price', s=10)
114 plt.title('Comparison of True Prices and Predicted Prices (2018-2020)')
115 plt.xlabel('Date')
116 plt.ylabel('Price')
117 plt.legend()
118 plt.grid(True)
119 plt.show()
```

```
[13765.01040549 14524.86439227 15116.96377777 ... 27231.21023167
27754.31058647 29188.28721168]
```

Result



Thank you!


```
most_recent_data = data[['Open', 'High', 'Low', 'Volume', 'Market Cap']].iloc[-1].values.reshape(1, -1)
```

```
next_day_prediction = random_forest_model.predict(most_recent_data)  
print("Predicted closing price for tomorrow:", next_day_prediction[0])
```

```
# =====
```

```
num_days = 7
```

```
input_features = np.array(data[['Open', 'High', 'Low', 'Volume', 'Market Cap']].iloc[-1]).reshape(1, -1)
```

```
predictions = []
```

```
for i in range(num_days):
```

```
    next_day_prediction = random_forest_model.predict(input_features)
```

```
    predictions.append(next_day_prediction[0])
```

```
    input_features = np.array([[next_day_prediction[0], next_day_prediction[0], next_day_prediction[0],  
                                input_features[0, 3], input_features[0, 4]]])
```

```
predictions
```

```
[43218.324227665114,  
 42708.869423206765,  
 42499.351815459064,  
 42152.798996691636,  
 41996.97053298736,  
 41930.95580938235,  
 41929.71168243981]
```

```
Predicted closing price for tomorrow: 43218.324227665114
```